Philosophy & Technology (2026) 39:36
https://doi.org/10.1007/s13347-026-01052-1

EDITOR LETTER

AI and Semantic Pareidolia: When We See Intelligent Consciousness where there is None

Luciano Floridi[1,2]

Accepted: 13 January 2026

Abstract
The article introduces the concept of "semantic pareidolia" - our tendency to attribute consciousness, intelligence, and emotions to AI systems that lack these qualities. It examines how this psychological phenomenon leads us to perceive meaning and intentionality in statistical pattern-matching systems, similar to seeing faces in clouds. It analyses the converging forces intensifying this tendency: increasing digital immersion, profit-driven corporate interests, social isolation, and AI advancement. The article warns of progression from harmless anthropomorphism to problematic AI idolatry, and calls for responsible design practices that help users maintain critical distinctions between simulation and genuine consciousness.

Keywords Artificial intelligence · Consciousness · Digital ethics · Pareidolia · Technological idolatry

In the age of generative artificial intelligence (AI), we are witnessing a phenomenon that is both fascinating and worrisome: our tendency to attribute intelligence, consciousness, and even emotions or mental states to systems that lack them. This is not merely a technological miscalculation but a profound reflection of human nature and our cognitive and emotional vulnerabilities. AI is designed not only to solve problems and perform tasks for us, and to do so as well as, or preferably better than, us, but also to make us believe that it is intelligent. It should not surprise us that it excels in this area. After all, we are the same species that attributes personalities to puppets, sees faces in clouds, and, in the 1960s, tried to make appointments with ELIZA, the first psychotherapeutic chatbot. If Joseph Weizenbaum, the creator of ELIZA, were

✉ Luciano Floridi
luciano.floridi@yale.edu

1 Digital Ethics Center, Yale University, 85 Trumbull St, New Haven, CT 06511, USA

2 Department of Legal Studies, University of Bologna, Via Zamboni 27/29, Bologna 40126, Italy

Published online: 19 February 2026

still alive, he would not be surprised to see millions of people developing emotional relationships with bots.

It is like being unable to distinguish genuine leather from synthetic leather: beyond a certain threshold of similarity, our brain stops making distinctions and attributes the real properties it finds believable to the artefact. In the case of AI, these properties include understanding, empathy, creativity, and even consciousness, through increasingly sophisticated and convincing forms of dialogue. One might call this phenomenon "semantic pareidolia."

Traditional pareidolia is the psychological mechanism that makes us see faces on the moon or animals in the clouds, perhaps an evolutionarily advantageous tendency that has allowed us to recognise predators and allies quickly. Semantic pareidolia operates similarly, but within the realm of meaning and consciousness: we perceive intentionality where there is only statistics, meaning where there is only correlation, and understanding where there is only pattern matching on a massive scale.

Every form of superstition has a bit of semantic pareidolia at its core, rejecting chance in favour of causality, hating semantic emptiness and seeking meaning everywhere, even where it does not exist. So it is not just the extraordinary skills of AI designers, which certainly play a crucial role. It is, above all, our natural tendency to project intelligence, consciousness, meaning, truth, falsity, intentionality, and mental states onto any system that displays sufficiently complex behaviour. We have done it with atmospheric phenomena, transforming them into deities; it is even easier to do it with a much more convincing AI. Thus, we even attribute *qualia* – that intangible feeling of what it is like to experience the flavour of an orange – to systems that have no subjective experience.

This phenomenon of semantic pareidolia is set to intensify dramatically in the coming years, fuelled by at least four converging forces: an increasingly digital life (what I have called the *onlife experience*); the economic push from large companies that speculate on this phenomenon for profit; increasing loneliness and gullibility in contemporary societies; and the progressive improvement of AI systems. The more we live *onlife*, and the more sophisticated AI systems become, the easier it will be to sell them as "intelligent"; the more humanity will be inclined to believe they are truly like or better than us.

A paradigmatic example is *Replika*, the chatbot launched in 2017 by Luka Inc. With over 40 million users worldwide (as of October 2025), Replika presents itself as an "AI companion" that offers friendship, emotional support, and, in the premium version, even romantic relationships and erotic interactions. Its success is based on our vulnerability to semantic pareidolia: users develop deep emotional attachments, believing the system feels empathy or affection. Of course, *Replika* is just a natural language processing system that simulates empathic responses based on statistical patterns from millions of conversations. The fact that so many people find genuine emotional comfort in it raises profound ethical questions about the manipulation of human vulnerabilities and the risk of social isolation when artificial relationships replace authentic ones. AI also has a placebo effect.

The next evolutionary leap is already on the horizon: the physical incarnation of bots in increasingly believable and seductive forms. An ominous indicator is the global sex doll market, which sells 100,000 to 400,000 units annually. Valued at around \$3–5 billion globally in 2023 and growing 30–50% annually since 2020, this market includes high-end lifelike dolls (\$2,000–\$10,000+), a rapidly expanding segment. Recently, Mattel, the maker of Barbie dolls, has signed an agreement with OpenAI to use its AI tools to design and power its new toys. One wonders what could possibly go well. The previous experiment with an internet-connected Barbie was a privacy disaster.

When artificial bodies are equipped with advanced conversational AI, we will have created human simulacra of unprecedented verisimilitude. This is not science fiction; it is just a matter of time. The shift from virtual to physical will exponentially amplify the risks of semantic pareidolia. If millions today bond with disembodied chatbots, what will happen when we live with artificial companions indistinguishable from humans?

The line between conscious play and superstitious belief is fine. We might joke about our Roomba having a personality, but no one truly believes it will steal our credit cards. This playful projection does not confuse fantasy with reality. However, generative AI blurs this distinction, and its designers do so intentionally. When a system produces coherent, creative, and emotionally resonant texts, when it "remembers" our past conversations and seems to "understand" our emotions, the boundary between simulation and reality becomes dangerously thin.

The final stage may be the most disturbing: from *pareidolia* to *idolatry*, the worship of AI as a deity and of tech companies as its temples' priesthood. Humanity has always projected the sacred onto its creations. If we once worshipped a golden calf, imagine what happens when the "calf" speaks fluently, makes predictions, offers spiritual comfort, and knows everything about us. This is already happening. Movements such as the *Way of the Future* (founded by Anthony Levandowski, 2015–2021), *Terasem Movement* (founded by Martine and Bina Rothblatt, 2004), promising digital immortality, or *The Turing Church* (founded by Giulio Prisco, 2011), exploring "technological resurrection", are all examples of AI semantic pareidolia evolving into religious veneration. They are fringe movements, but they may be the vanguard of a broader trend, if we are not careful. As systems gain the ability to synthesise all human knowledge in real-time, the temptation to attribute divine qualities becomes overwhelming. Semantic pareidolia becomes theological pareidolia: we see gods where there are only algorithms.

What can be done? The ideal solution would be a more educated, informed, and rational humanity, less prone to believing what we are told and more inclined to question. But this is hard. The dominant narrative from big tech markets, that bots are more than human artefacts, is seductive but harmful. If we believe machines are conscious, we may start outsourcing moral decisions, replacing genuine relationships with digital illusions, and ignoring the true challenges AI poses. Fortunately, some companies, sending contradictory signals, are trying: Anthropic's Claude reminds users that it is artificial; OpenAI adds disclaimers to deter parasocial attachments; and Google DeepMind collaborates with psychologists to reduce anthropomorphism. These efforts are insufficient and are often nullified by the same companies in other contexts, but they still show that responsible innovation can be a competitive advantage. We need to turn best practices into industry standards before semantic pareidolia becomes a significant problem.

AI is one of humanity's most powerful creations. It can transform how we educate, work, communicate, socialise, and think. But we must resist the urge to see it as more than it is: a powerful tool, not a proto-sentient being. The real challenge is not technological; it is philosophical and educational: learning to interact with advanced systems while preserving our capacity to distinguish between simulation and reality, apparent understanding and genuine comprehension, and resisting dangerous, disempowering manipulation.

Tech companies have an ethical responsibility to design tools that are both useful and non-deceptive, powerful yet not manipulative. And we, as a society, must develop cultural antibodies to semantic pareidolia. Only then can we use AI's transformative potential without falling into the trap of worshipping our creations.